

Saraswathi (Kunju) Menon

Flower Mound, TX 75022 | 469-386-3470 | sarakmenon@gmail.com | [Linkedin](#) | [Website](#)

EDUCATION

The University of Texas at Dallas

Major in Computer Science, Bachelor of Science

Expected Graduation Date: May 2027

GPA 4.0

SKILLS

Programming Languages: Python, Java, C, C++, SQL, JavaScript

Frameworks/Libraries: LangChain, LangGraph, GPT-4o, RAG, LLMs, Prompt Engineering, React, Node.js, FastAPI, Firestore, PostgreSQL

Relevant Courses: Data Structures & Algorithms, Software Engineering, Operating Systems, Computer Architecture, UNIX/Linux

Tools: Git (Version Control), Azure DevOps, Terraform, MCP, OAuth2/JWT, Docker, Kubernetes, Copilot, Codex, Windsurf, CI/CD, AWS

PROFESSIONAL EXPERIENCE

Toshiba Global Commerce Solutions, Frisco, TX

May 2026 - Present

Agentic Software Engineering Intern | Intern Ambassador

- Developed an AI-powered Azure deployment automation platform using agentic workflows, specialized AI workers, and LLM reasoning to orchestrate Terraform infrastructure and production deployment workflows, increasing enterprise deployment automation by 40% and improving operational decision-making.
- Developed production-grade AI microservices and agentic workflows integrating SaaS APIs, deployment metadata, and REST services to automate enterprise workflows and reduce manual input correction by 20%.
- Collaborated with cross-functional engineers to build backend orchestration and control-plane integrations for multi-agent deployment workflows, improving reliability and routing accuracy by 30%.

Amdocs @AT&T, Plano, TX

Aug 2025 – Dec 2025

Generative AI Engineer (Part-Time)

NetBond Advanced Chat

- Developed an enterprise agentic AI platform using LangGraph Studio, Model Context Protocol (MCP), and Transformer Models to orchestrate AI agents and automate secure cloud configuration generation.
- Engineered modular RAG pipelines integrating MCP servers, Python FastAPI REST APIs, and enterprise knowledge sources to enable AI agents to retrieve domain context data and execute external tools for accurate cloud configuration recommendations.
- Designed an end-to-end AI application combining LLM reasoning, prompt engineering, and a React/Node.js interface to enable real-time generation, validation, and editing of cloud configurations.

AskVoice

- Migrated and refactored Python FastAPI REST APIs and MongoDB data integrations to updated service architecture, ensuring compatibility and improved performance across deployment environments.
- Deployed applications through Azure DevOps CI/CD pipelines, supporting integration testing and production validation while using KQL and Azure Data Explorer (ADX) logs for debugging and troubleshooting system issues and accelerate root-cause analysis.

MATH-101  Remote

Aug 2025 – Present

Full-Stack Web Developer

- Developed and deployed a responsive, consumer-facing web application using React, Next.js, TypeScript, and Tailwind CSS, building reusable UI design components and user workflows for students and administrators, supporting 7 math courses.
- Engineered full-stack application features including protected routes, role-based access control, account approval, course assignment, automated email workflows, and an administrative dashboard integrated with Firebase backend services.
- Improved application reliability, security, and user experience through responsive UI development, input validation, rate limiting, server-side token verification, Firebase Security Rules, automated testing, secure coding, and defensive programming practices.

Amdocs @AT&T, Plano, TX

June 2025 - Aug 2025

Generative AI / Data Science Intern

- Developed cloud-agnostic agentic AI workflows using LangGraph, LangChain, GPT-4o, and prompt engineering, incorporating domain-specific logic, semantic similarity scoring, and LLM-based reasoning to automate enterprise marketing content generation while ensuring adherence to AT&T creative and compliance standards.
- Developed production-ready FastAPI AI services with OAuth2/JWT authentication and Retrieval-Augmented Generation (RAG), enabling secure enterprise document search through natural-language APIs and AI-powered retrieval.
- Designed scalable AI data pipelines that transformed unstructured PDFs and PowerPoint documents into structured JSON datasets for downstream retrieval, analysis, and automated content generation, reducing document-processing time by 30%.

PROJECTS

OptiClarity - Deep Learning Computer Vision Diagnostic System

Jan 2025 – May 2025

- Developed an AI-powered eye health platform using PyTorch, TensorFlow, and ResNet50 CNN architectures to analyze four types of retinal scans across four deep learning models, including an ensemble model for multi-modal disease classification.
- Engineered an end-to-end AI diagnostic workflow with a LangChain-based RAG assistant, enabling image-based disease prediction and context-aware eye-health guidance through a React and Python web application.

CERTIFICATIONS

- AWS Certified AI Practitioner (AIF-C01) — Amazon Web Services (AWS) | Aug 2026